

AthenaPK

Philipp Grete
University of Hamburg

in (development) collaboration with
Forrest Glines and James Stone

Parthenon developer meeting



Universität Hamburg
DER FORSCHUNG | DER LEHRE | DER BILDUNG



Funded by the
European Union


AthenaPK – Overview

Applications (AthenaPK)

- Finite-volume (M)HD
- Cell-centered divergence cleaning

Parthenon interface

- Single package with MultiStageDriver
- MeshData where supported
- One large variable vector
- Explicit ScratchPadMemory use
- C++17 constexpr if in kernels
- Regression test infrastructure
- Lots of Params abuse




Application: Turbulence

Background

- 3D periodic, uniform grids
- May include cooling (optically thin) and diffusion processes
- Study energy dynamics \Rightarrow requires high resolution

Implementation

- Turbulence driver (stochastic process in spectral space for force field in physical space)
- Custom, explicit Fourier transform using few modes



- Data compression
- Arbitrary “MeshData” (say like an array of acceleration vectors)
- Support for slices in outputs

Application: Cloud in wind

Background

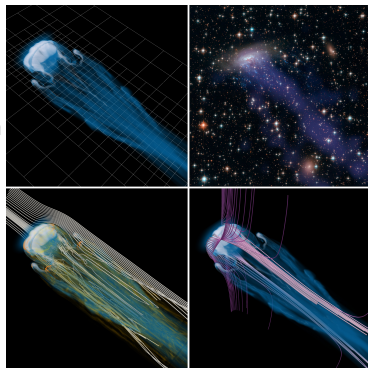
- Supersonic ($M \approx 4$), weak magnetic fields ($\beta = 100$), high density contrast ($\rho_c/\rho_w = 1000$), with AMR
- Includes cooling, anisotropic thermal conduction
- Study cloud crushing/survival

Implementation

- RKL2 super timestepping (separate “driver” \rightarrow just more tasks)
- first-order flux correction (iteratively, not guaranteed to work)

Parthenon wish list

- “fix” prolong/restriction
- Face-centered fields (for CT)



- FluxDiv subset of variables
- boundary exchange subset of variables

Application: AGN feedback

Background


- SMR with static potential
- Single AGN in the center
- All the previous physics, too
- Study AGN feeding, energy redistribution, magnetization, ...

Implementation

- global reduction/kernel

Parthenon wish list

- adaptive time-stepping?
- use texture cache for cooling table (and interpolation)



Application: “HPC”

Background

- used in CFP benchmarking suite for upcoming cluster in Hamburg
- used in acceptance testing for Frontier

Parthenon wish list

- further opt./tests on AMD GPUs (and Intel)
- HostPinned versus Device memory
- optimized buffer filling kernels dep. on number of buffers
- “auto-tuning”